

*#Using sklearn.datasets.load\_diabetes apply Variance method for removing the constant column also after applying  
#the Variance method apply multi linear regression on that data  
/aleternate dataset given by ajay sharma and told to apply svm on it*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.cluster import DBSCAN
from sklearn import metrics
from sklearn.datasets import make_blobs
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt

from sklearn.feature_selection import VarianceThreshold

df=pd.read_csv("E:\diabetesnew.csv")
var_thres=VarianceThreshold(threshold=0.2)
var_thres.fit(df)
var_thres.get_support()
df.columns[var_thres.get_support() == True]
columns_having_var_more_than_50 = df.columns[var_thres.get_support()
== True]

columns_having_var_less_than_50 = df.columns[var_thres.get_support()
== False]
df.drop(columns_having_var_less_than_50,inplace = True,axis= 1)
df.isnull().values.any()

False

from sklearn import svm
from sklearn.model_selection import train_test_split

x = df.iloc[:, :-2]
y = df.iloc[:, -1]
x_train, x_test, y_train, y_test = train_test_split(x, y, random_state
= 0, test_size = 0.2)

clf = svm.SVC(kernel='rbf')
clf.fit(x_train,y_train)
y_pred = clf.predict(x_test)

from sklearn.metrics import accuracy_score
print("Accuracy:", accuracy_score(y_test, y_pred))

Accuracy: 0.7987012987012987
```

*#Using sklearn.datasets.load\_wine Apply Correlation and make a heat map using seaborn and remove the highly  
#correlated columns if exist and the apply SVM and get the best accuracy by changing the Hyperparameters*

```

import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
%matplotlib inline

from sklearn.datasets import load_wine
wine=load_wine()
type(wine)
print(wine.keys())

dict_keys(['data', 'target', 'frame', 'target_names', 'DESCR',
'feature_names'])

columns_name=wine.feature_names

df = pd.DataFrame(wine['data'],columns=wine['feature_names'])

df["medv"]=wine.target


x=df.drop("medv",axis=1)
y=df["medv"]

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=40)
x_train.shape,x_test.shape

((124, 13), (54, 13))

import seaborn as sns
plt.figure(figsize=(12,10))
cor=x_train.corr()
sns.heatmap(cor,annot=True)
plt.show()

```



```
def correlation(dataset, threshold):#X_train,
    col_corr = set() # Set of all the names of correlated columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix)): #traverse through the rows
        for j in range(i): #traverse through column
            if abs(corr_matrix.iloc[i, j]) > threshold: # we are
interested in absolute coeff value
                colname = corr_matrix.columns[i] # getting the name
of column
                col_corr.add(colname)
    return col_corr

corr_features = correlation(x_train, 0.3)
len((corr_features))

10

df.head()

    alcohol  malic_acid  ash  alcalinity_of_ash  magnesium
total_phenols  \
```

0	14.23	1.71	2.43	15.6	127.0
2.80					
1	13.20	1.78	2.14	11.2	100.0
2.65					
2	13.16	2.36	2.67	18.6	101.0
2.80					
3	14.37	1.95	2.50	16.8	113.0
3.85					
4	13.24	2.59	2.87	21.0	118.0
2.80					

	flavanoids	nonflavanoid_phenols	proanthocyanins	color_intensity
hue \				
0	3.06	0.28	2.29	5.64
1.04				
1	2.76	0.26	1.28	4.38
1.05				
2	3.24	0.30	2.81	5.68
1.03				
3	3.49	0.24	2.18	7.80
0.86				
4	2.69	0.39	1.82	4.32
1.04				

	od280/od315_of_diluted_wines	proline	medv
0	3.92	1065.0	0
1	3.40	1050.0	0
2	3.17	1185.0	0
3	3.45	1480.0	0
4	2.93	735.0	0

```
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
```

```
feature_df=df[["alcohol","malic_acid","ash"]]
x=np.asarray(feature_df)
x[0:5]
```

```
array([[14.23,  1.71,  2.43],
       [13.2 ,  1.78,  2.14],
       [13.16,  2.36,  2.67],
       [14.37,  1.95,  2.5 ],
       [13.24,  2.59,  2.87]])
```

```
wine['medv']=df['medv'].astype('int')
y=np.asarray(wine['medv'])
print("y values:",y[0:5])
```

```
y values: [0 0 0 0 0]
```

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.2,random
_state=50)
print("Train set:",x_train.shape,y_train.shape)
print("Test set:",x_test.shape,y_test.shape)
```

```
Train set: (142, 3) (142,)
Test set: (36, 3) (36,)
```

```
from sklearn import svm
clf=svm.SVC(kernel='poly')
clf.fit(x_train,y_train)
```

```
yhat=clf.predict(x_test)
yhat[0:5]
```

```
array([1, 1, 1, 2, 2])
```

```
from sklearn.metrics import f1_score
f1_score(y_test,yhat,average='weighted')
```

```
0.75
```

```
clf2=svm.SVC(kernel='rbf')
clf2.fit(x_train,y_train)
yhat2=clf2.predict(x_test)
print("avg f1-score:%.4f"% f1_score(y_test,yhat2,average='weighted'))
```

```
avg f1-score:0.7467
```

```
clf2=svm.SVC(kernel='linear')
clf2.fit(x_train,y_train)
yhat2=clf2.predict(x_test)
print("avg f1-score:%.4f"% f1_score(y_test,yhat2,average='weighted'))
```

```
avg f1-score:0.7235
```

*#Using sklearn.datasets.load\_diabetes apply Mutual info Classification and check which are the best columns*

*#according to the target column.*

*#Then Apply decision tree on that data and try to get best accuracy by changing the hyperparameters*

```
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import numpy as np
%matplotlib inline
```

```
from sklearn.datasets import load_diabetes
data =load_diabetes ()
type(data)
```

```

sklearn.utils.Bunch
data.keys
<function Bunch.keys>
data.feature_names
['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']
columns_name = data.feature_names
columns_name
['age', 'sex', 'bmi', 'bp', 's1', 's2', 's3', 's4', 's5', 's6']
df = pd.DataFrame(data.data, columns = columns_name)
df.head()

```

	age	sex	bmi	bp	s1	s2
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163
2	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596

	s4	s5	s6
0	-0.002592	0.019908	-0.017646
1	-0.039493	-0.068330	-0.092204
2	-0.002592	0.002864	-0.025930
3	0.034309	0.022692	-0.009362
4	-0.002592	-0.031991	-0.046641

```

df["daibetes"] = data.target #dependent columnn CONTENT
X = df.drop("daibetes",axis=1) #independent variable : all column except Target Dv colun
y = df["daibetes"] #dependent variables only target column will be in Y
X_train,X_test,y_train,y_test=train_test_split(X, #INDEPENDENT VARIABLE
y, # as DEPENDENT VARIABLE
test_size=0.3, #70% TRAINING DS AND 30% TEST DATA
random_state=0)

```

```
X_train.shape
```

```
(309, 10)
```

```
from sklearn.feature_selection import mutual_info_regression
```

```
# determine the mutual information
```

```
mutual_info = mutual_info_regression(X_train, y_train)
```

```
mutual_info #impactful variable will get high value and less  
impactfull will get low values
```

```
array([0.01410304, 0.03507104, 0.19056747, 0.10522933, 0.08487774,  
       0.00482829, 0.05602829, 0.12500372, 0.15324809, 0.1194692 ])
```

```
mutual_info = pd.Series(mutual_info)
```

```
mutual_info.index = X_train.columns
```

```
mutual_info.sort_values(ascending=False)
```

```
bmi    0.190567  
s5     0.153248  
s4     0.125004  
s6     0.119469  
bp     0.105229  
s1     0.084878  
s3     0.056028  
sex    0.035071  
age    0.014103  
s2     0.004828  
dtype: float64
```

```
from sklearn.feature_selection import SelectKBest
```

```
from sklearn.feature_selection import mutual_info_regression
```

```
sel_best_cols = SelectKBest(mutual_info_regression, k=5)
```

```
sel_best_cols.fit(X_train, y_train)
```

```
SelectKBest(k=5,  
            score_func=<function mutual_info_regression at  
0x00000247DDAC9D30>)
```

```
X_train.columns[sel_best_cols.get_support()==True]
```

```
Index(['bmi', 'bp', 's4', 's5', 's6'], dtype='object')
```

```
type(X_train)
```

```
pandas.core.frame.DataFrame
```

```
X_train = X_train[['sex', 'bmi', 's3', 's4', 's5']]
```

```
X_test = X_test[['sex', 'bmi', 's3', 's4', 's5']]
```

```

from sklearn.tree import DecisionTreeRegressor
tree=DecisionTreeRegressor()(criterion="entropy",max_depth=4)
tree

DecisionTreeRegressor()

tree.fit(X_train,y_train)

DecisionTreeRegressor()

y_pred=tree.predict(X_test)
#y_pred=tree.predict(X_test)

print(y_pred[0:5])
print(y_test[0:5])

[261. 310. 225. 214. 191.]
362      321.0
249      215.0
271      127.0
435       64.0
400      175.0
Name: daibetes, dtype: float64

from sklearn import metrics
print("DecisionTrees's Accuracy: ",metrics.r2_score(y_pred,y_test))

DecisionTrees's Accuracy:  -0.149140648377031

from sklearn.metrics import mean_squared_error

rmse = (np.sqrt(mean_squared_error(y_pred,y_test)))
rmse

83.59538246772253

```

```

#Using sklearn.datasets.load_boston apply Mutual info Regression and
check which are the best columns according
#to the target column.
#Then Apply MultiLinear Regression on that data and try to get best
accuracy by changing the hyperparameters
from sklearn.datasets import load_boston
data =load_boston ()
type(data )

```

```

C:\Users\Renuka\anaconda3\lib\site-packages\sklearn\utils\
deprecation.py:87: FutureWarning: Function load_boston is deprecated;
`load_boston` is deprecated in 1.0 and will be removed in 1.2.

```

The Boston housing prices dataset has an ethical problem. You can refer to



the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np

data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22,
header=None)
data = np.hstack([raw_df.values[::2, :],
raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)

sklearn.utils.Bunch
data.feature_names
array(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
'RAD',
'TAX', 'PTRATIO', 'B', 'LSTAT'], dtype='<U7')

import pandas as pd
columns_name = data.feature_names
```

```
df = pd.DataFrame(data.data, columns = columns_name)
df.head()
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

```
df["dv"] = data.target #dependent columnn CONTENT
```

```
X = df.drop("dv",axis=1) #independent variable : all column except Target Dv colun
```

```
y = df["dv"] #dependent variables only target column will be in Y
```

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(X, #INDEPENDENT VARIABLE
        y, # as DEPENDENT VARIABLE
        test_size=0.3, #70% TRAINING DS AND 30% TEST DATA
        random_state=0)
```

```
X_train.shape
```

```
(354, 13)
```

```
y_train.shape
```

```
(354,)
```

```
from sklearn.feature_selection import mutual_info_regression
# determine the mutual information
mutual_info = mutual_info_regression(X_train, y_train)
mutual_info #impactful variable will get high value and less impactfull will get low values
```

```
array([0.32559106, 0.18773109, 0.53440947, 0.03137659, 0.43582696,  
       0.59739155, 0.33866851, 0.30723383, 0.22056715, 0.36521941,  
       0.51343429, 0.16564181, 0.65114809])
```

```
X_train.shape
```

```
(354, 13)
```

```
mutual_info = pd.Series(mutual_info)  
mutual_info.index = X_train.columns  
mutual_info.sort_values(ascending=False)
```

```
LSTAT      0.651148  
RM         0.597392  
INDUS      0.534409  
PTRATIO    0.513434  
NOX        0.435827  
TAX        0.365219  
AGE        0.338669  
CRIM       0.325591  
DIS        0.307234  
RAD        0.220567  
ZN         0.187731  
B          0.165642  
CHAS       0.031377  
dtype: float64
```

```
from sklearn.feature_selection import SelectKBest  
from sklearn.feature_selection import mutual_info_regression
```

```
sel_best_cols = SelectKBest(mutual_info_regression, k=5)
```

```
sel_best_cols.fit(X_train, y_train)
```

```
SelectKBest(k=5,  
            score_func=<function mutual_info_regression at  
0x00000247DDAC9D30>)
```

```
X_train.columns[sel_best_cols.get_support()==True]
```

```
Index(['INDUS', 'NOX', 'RM', 'PTRATIO', 'LSTAT'], dtype='object')
```

```
X_train = X_train[['INDUS', 'NOX', 'RM', 'PTRATIO', 'LSTAT']]
```

```
X_train.shape
```

```
(354, 5)
```

```
X_test = X_test[['INDUS', 'NOX', 'RM', 'PTRATIO', 'LSTAT']]
```

```
y_train.shape
```

```
(354,)
```

```

X_train.shape
(354, 5)
X_train = X_train.values

import numpy as np
from sklearn import linear_model
regr = linear_model.LinearRegression()
regr.fit(X_train,y_train)#training func question + answers
# The coefficients
print ('Intercept: ',regr.intercept_)
print ('Coefficient : ',regr.coef_)

Intercept: 24.661233547043977
Coefficient : [ 0.0533226 -6.52551265  4.57081317 -1.15326995 -
0.51504091]

from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error

y_hat = regr.predict(X_train)
#rmse = (np.sqrt(mean_squared_error(y_train, y_hat)))

r2 = r2_score(y_train, y_hat)

r2
0.7123290285129122

y_test_predict = regr.predict(X_test)
#rmse = (np.sqrt(mean_squared_error(y_test, y_test_predict)))
r2 = r2_score(y_test,y_test_predict)

C:\Users\Renuka\anaconda3\lib\site-packages\sklearn\base.py:443:
UserWarning: X has feature names, but LinearRegression was fitted
without feature names
  warnings.warn(

r2
0.5948238037827576

```